



EE 542

Lecture 18: Spark on Cloud Internet and Cloud Computing

Young Cho

Department of Electrical Engineering

University of Southern California

MapReduce

- Programming model for data-intensive computing on commodity clusters
- Pioneered by Google
 - Processes 20 PB of data per day
- Popularized by Apache Hadoop project
 - Used by Yahoo!, Facebook, Amazon, ...
- Industry Trending Toward Spark
 - Quickly Being Adopted by Big Data Industry

Limitations of MapReduce

- MapReduce is great at one-pass computation, but inefficient for *multi-pass* algorithms
- No efficient primitives for data sharing
 - State between steps goes to distributed file system
 - Slow due to replication & disk storage
 - No control of data partitioning across steps

Spark Programming Model

- Extends MapReduce with primitives for efficient data sharing
 - “Resilient distributed datasets”
- Open source in Apache Incubator
 - Growing community with 100+ contributors
- APIs in Java, Scala & Python

Resilient Distributed Datasets (RDDs)


- Collections of objects stored across a cluster
- User-controlled partitioning & storage (RAM, disk, ...)
- Automatically rebuilt on failure

```
urls = spark.textFile("hdfs://...")
records = urls.map(lambda s: (s, 1))
counts = records.reduceByKey(lambda a, b: a + b)
bigCounts = counts.filter(lambda (url, cnt): cnt > 10)
```

```
bigCounts.cache()
```

```
bigCounts.filter(
    lambda (k,v): "news" in k).count()
```

```
bigCounts.join(otherPartitionedRDD)
```



Spark


- In-Memory Computation
 - For 64-bit computers TB of data in RAM
 - Designed to transform data in-mem and not in disk
 - Supports parallel distributed processing of data
 - 100x in memory and 10x on disk then Hadoop
- General programming model
 - Use normal sequential programming
 - No need for maps and reduce operations

Spark: Key Advantages

- Ability for On-disk Data Sorting
 - Tuned for large scale of data sorting on disk
 - The world record of on-disk large scale data sorting
- Efficient Use of Cache
 - Mesos which is a distributed system kernel for caching the intermediate dataset
 - Multiple iterations on the cached dataset
- In-memory Tuned Library
 - MLlib library for in-memory tuned operations
- Faster Launch with Virtual Machine


Clustering

Grouping **data** according to similarity



Clustering

Grouping data according to similarity



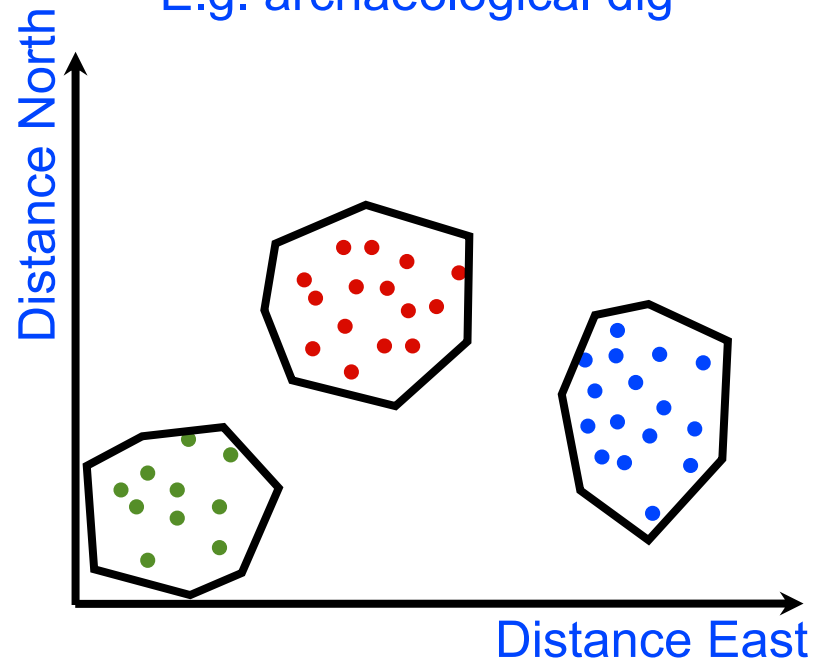
Clustering with Spark

K-Means: preliminaries

Benefits

- Popular
- Fast
- Conceptually straightforward

E.g. archaeological dig




Clustering with Spark

K-Means: preliminaries

Data: Collection of values

```
data = lines.map(line=>  
  parseVector(line))
```




Clustering with Spark

K-Means: preliminaries

Dissimilarity:

Squared Euclidean distance

```
dist = p.squaredDist(q)
```



Clustering with Spark


K-Means: preliminaries

K = Number of clusters

$\mu_1, \mu_2, \dots, \mu_K$

Data assignments to clusters

S_1, S_2, \dots, S_K



Clustering with Spark


K-Means: preliminaries

K = Number of clusters

$\mu_1, \mu_2, \dots, \mu_K$

Data assignments to clusters


S_1, S_2, \dots, S_K



Clustering with Spark

K-Means Algorithm


- Initialize K cluster centers
- Repeat until convergence:
 - Assign each data point to the cluster with the closest center.
 - Assign each cluster center to be the mean of its cluster's data points.



Clustering with Spark

K-Means Algorithm

- Initialize K cluster centers
- Repeat until convergence:
 - Assign each data point to the cluster with the closest center.
 - Assign each cluster center to be the mean of its cluster's data points.




Clustering with Spark

K-Means Algorithm

- Initialize K cluster centers

```
centers = data.takeSample(  
    false, K, seed)
```

- Repeat until convergence:
 Assign each data point to
 the cluster with the
 closest center.
 Assign each cluster
 center to be the mean of
 its cluster's data points.




Clustering with Spark

K-Means Algorithm

- Initialize K cluster centers

```
centers = data.takeSample(  
    false, K, seed)
```
- Repeat until convergence:
 - Assign each data point to the cluster with the closest center.
 - Assign each cluster center to be the mean of its cluster's data points.




Clustering with Spark

K-Means Algorithm

- Initialize K cluster centers

```
centers = data.takeSample(  
    false, K, seed)
```
- Repeat until convergence:
 Assign each data point to
 the cluster with the
 closest center.
 Assign each cluster
 center to be the mean of
 its cluster's data points.



Clustering with Spark


K-Means Algorithm

- Initialize K cluster centers

```
centers = data.takeSample(  
    false, K, seed)
```

- Repeat until convergence:

```
closest = data.map(p =>  
  
(closestPoint(p,centers),p))  
Assign each cluster  
center to be the mean of  
its cluster's data points.
```



Clustering with Spark

K-Means Algorithm

- Initialize K cluster centers


```
centers = data.takeSample(  
    false, K, seed)
```

- Repeat until convergence:

```
closest = data.map(p =>
```

```
(closestPoint(p,centers),p))
```

Assign each cluster
center to be the mean of
its cluster's data points.



Clustering with Spark

K-Means Algorithm

- Initialize K cluster centers


```
centers = data.takeSample(  
    false, K, seed)
```

- Repeat until convergence:

```
closest = data.map(p =>
```

```
(closestPoint(p,centers),p))
```

Assign each cluster
center to be the mean of
its cluster's data points.



Clustering with Spark


K-Means Algorithm

- Initialize K cluster centers

```
centers = data.takeSample(  
    false, K, seed)
```

- Repeat until convergence:

```
closest = data.map(p =>  
  
    (closestPoint(p,centers),p))  
pointsGroup =  
    closest.groupByKey()
```



Clustering with Spark

K-Means Algorithm

- Initialize K cluster centers

```
centers = data.takeSample(  
    false, K, seed)
```


- Repeat until convergence:

```
closest = data.map(p =>
```

```
(closestPoint(p, centers), p))
```

```
pointsGroup =  
    closest.groupByKey()
```

```
newCenters = pointsGroup.mapValues(  
    ps => average(ps))
```



Clustering with Spark

K-Means Algorithm

- Initialize K cluster centers

```
centers = data.takeSample(  
    false, K, seed)
```


- Repeat until convergence:

```
closest = data.map(p =>
```

```
(closestPoint(p, centers), p))
```

```
pointsGroup =  
    closest.groupByKey()
```

```
newCenters = pointsGroup.mapValues(  
    ps => average(ps))
```



Clustering with Spark


K-Means Algorithm

- Initialize K cluster centers

```
centers = data.takeSample(  
    false, K, seed)
```

- Repeat until convergence:

```
closest = data.map(p =>  
    (closestPoint(p,centers),p))  
pointsGroup =  
    closest.groupByKey()  
newCenters = pointsGroup.mapValues(  
    ps => average(ps))
```



Clustering with Spark


K-Means Algorithm

- Initialize K cluster centers

```
centers = data.takeSample(  
    false, K, seed)
```

- Repeat until convergence:

```
while (dist(centers,  
            newCenters) >  $\epsilon$ )  
  
closest = data.map(p =>  
    (closestPoint(p, centers), p))  
pointsGroup =  
    closest.groupByKey()  
newCenters = pointsGroup.mapValues(  
    ps => average(ps))
```



Clustering with Spark


K-Means Algorithm

- Initialize K cluster centers

```
centers = data.takeSample(  
    false, K, seed)
```

- Repeat until convergence:


```
while (dist(centers,  
            newCenters) >  $\epsilon$ )  
    closest = data.map(p =>  
        (closestPoint(p, centers), p))  
    pointsGroup =  
        closest.groupByKey()  
    newCenters = pointsGroup.mapValues(  
        ps => average(ps))
```



Clustering with Spark

K-Means Source

```
centers = data.takeSample(
  false, K, seed)
while (d >  $\epsilon$ )
{
  closest = data.map(p =>
    (closestPoint(p, centers), p))
  pointsGroup =
    closest.groupByKey()
  newCenters = pointsGroup.mapValues(
    ps => average(ps))
  d = distance(centers, newCenters)
  centers = newCenters.map(_)
}
```




Clustering with Spark

Ease of use


- Interactive shell:
 - Useful for featurization, pre-processing data
- Lines of code for K-Means
 - Spark ~ 90 lines
 - Hadoop/Mahout ~ 4 files, > 300 lines

Performance

K-Means



Logistic Regression



[Zaharia et. al, NSDI'12]

K-Means in MLlib

- <http://spark.apache.org/docs/latest/mllib-clustering.html#k-means>
- Available for Multiple Languages
 - Scala
 - Java
 - Python