



EE 542

Lecture 8: Single Processor to Cloud


Internet and Cloud Computing

Young Cho

Department of Electrical Engineering

University of Southern California


Moore's Law



Data source: Wikipedia (https://en.wikipedia.org/wiki/Transistor_count)

The data visualization is available at OurWorldinData.org. There you find more visualizations and research on this topic.

CPU Scaling



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2017 by K. Rupp

<https://www.karlrupp.net/2018/02/42-years-of-microprocessor-trend-data/>

Parallelism

- Bit Level
 - Parallel Circuit Design
- Instruction Level
 - Microarchitecture
 - Example: Pipelining and Superscalar
- Thread/Task Level
 - Independent Processes
 - Multiple cores, register files, and shared memory
- Application Level
 - Independent Programs
 - Multiple cores and Multiple chips

Instruction-Level Parallelism (ILP)

- Parallelism at the machine-instruction level
- The processor can re-order, pipeline instructions, split them into microinstructions, do aggressive branch prediction, etc.
- Instruction-level parallelism enabled rapid increases in processor speeds

Thread-Level Parallelism (TLP)

- This is parallelism on a more coarser scale
- Server can serve each client in a separate thread (Web server, database server)
- A computer game can do AI, graphics, and physics in three separate threads
- Single-core superscalar processors cannot fully exploit TLP
- Multi-core architectures are the next step in processor evolution: explicitly exploiting TLP

Application Level Parallelism

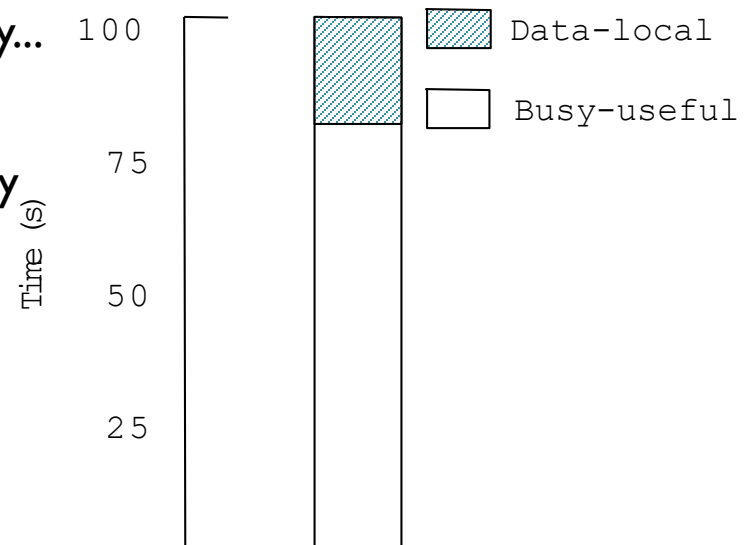

- Multicore Processor
 - All processors are on the same chip
 - Multi-core processors are MIMD
 - Multiple Instructions Multiple Data
 - Extensive Shared Memory Management
- Clustered Computers over Network
 - Scales based on the number of parallel programs
 - Potential for tighter integration - difficult to achieve

The Ultimate Goal


- Infinite Scalability
 - More HW => Higher Performance
- Difficulty in Parallelization
 - Need to parallelize a single thread into independent threads
 - Limits dependent on the algorithm
 - Or is there a limit??

Uniprocessor View

- Performance depends heavily on memory hierarchy
- Time spent by a program
 - $\text{Timeprog}(I) = \text{Busy}(I) + \text{Data Access}(I)$
 - Divide by cycles to get CPI equation
- Data access time can be reduced by:
 - Optimizing machine
 - bigger caches, lower latency...
 - Optimizing program
 - temporal and spatial locality




Uniprocessor vs. Multiprocessor




Multiprocessor

- A collection of communicating processors
 - Goals: balance load, reduce inherent communication and extra work



- A multi-cache, multi-memory system
 - Role of these components essential regardless of programming model
 - Prog. model and comm. abstr. affect specific performance tradeoffs



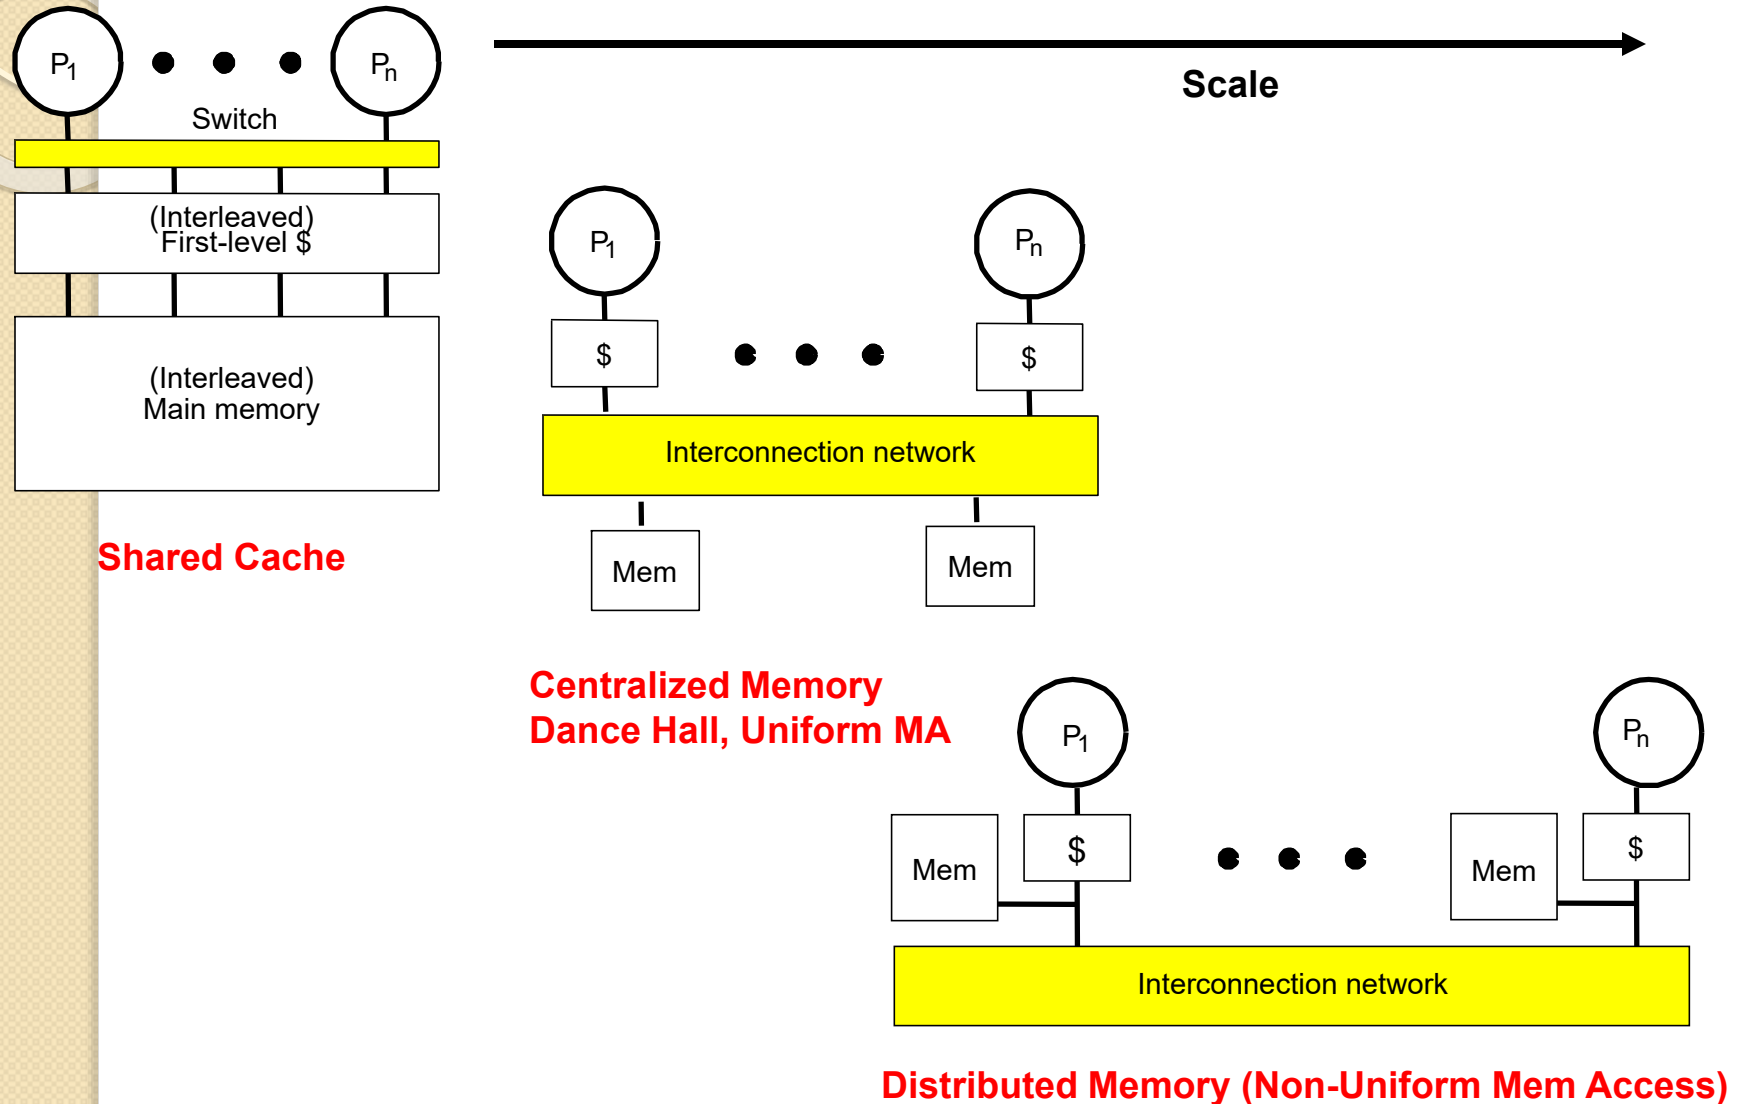
Parallel Architecture

- Parallel Architecture
 - Computer Architecture
 - Communication Architecture
- Small-scale shared memory
 - Extend the memory system to support multiple processors
 - Good for multiprogramming throughput and parallel computing
 - Allows *fine-grain sharing* of resources
- Characterization
 - Naming
 - Synchronization
 - Performance

Characterization


- Naming
 - Global physical address space
 - Global virtual address space
 - Segmented shared address space
- Synchronization
 - Local and Shared Memory
 - BUS and Network Latencies
- Performance of Parallel Algorithms
 - Total Latency
 - Total Bandwidth
 - Often not the processor

Memory Configurations




Shared Cache

- Alliant FX-8
 - early 80's
 - eight 68020s with x-bar to 512 KB interleaved cache
- Encore & Sequent
 - first 32-bit micros (N32032)
 - two to a board with a shared cache




Advantages

- Single Cache
 - Only one copy of any cached block
- Fine-Grain Sharing
 - Cray Xmp has shared registers!
- Potential for Positive Interference
 - One proc pre-fetches data for another
- Smaller Total Storage
 - Only one copy of code/data
- Can share data within a line
 - Long lines without false sharing




Disadvantages

- Fundamental BW limitation
- Increases latency of all accesses
 - Cross-bar (Multiple ports)
 - Larger cache
 - LI hit time determines proc. Cycle
- Potential for negative interference
 - one proc flushes data needed by another
- Many L2 caches are shared today
- CMP makes cache sharing attractive



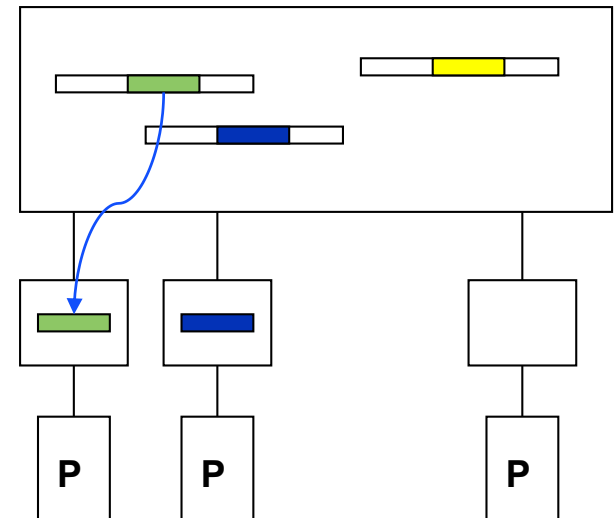
Bus-Based Symmetric Shared Memory



- Dominate the server market
 - Building blocks for larger systems; arriving to desktop
- Attractive as throughput servers and for parallel programs
 - Fine-grain resource sharing
 - Uniform access via loads/stores
 - Automatic data movement and coherent replication in caches
 - Cheap and powerful extension
- Normal uniprocessor mechanisms to access data
 - Key is extension of memory hierarchy to support multiple processors
- Chip Multiprocessors

Caches are Critical for Performance

- Reduce average latency
 - automatic replication closer to processor
- Reduce average bandwidth
- Data is logically transferred from producer to consumer to memory
 - store reg --> mem
 - load reg <-- mem
- Many processors can shared data efficiently
- What happens when store & load are executed on different processors?



Cache Coherence

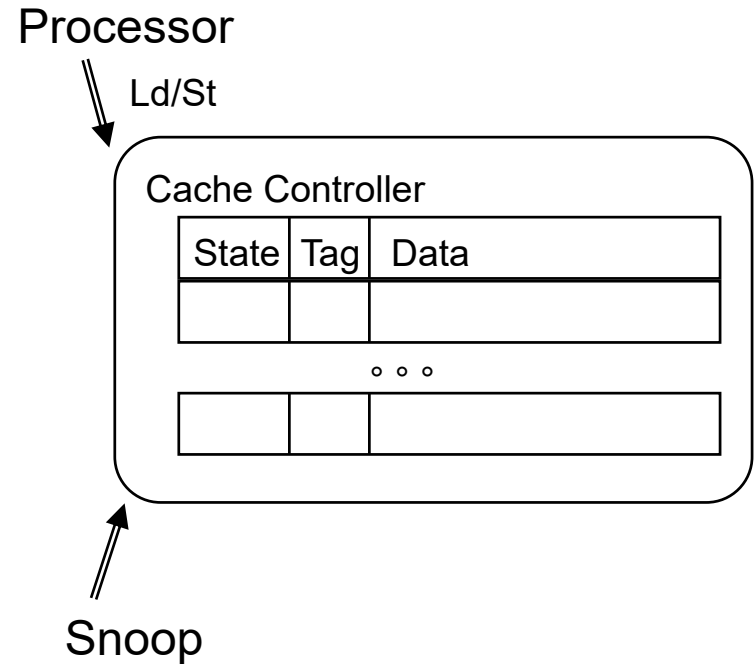
- Caches play key role in all cases
 - Reduce average data access time
 - Reduce bandwidth demands placed on shared interconnect
- Private processor caches create a problem
 - Copies of a variable can be present in multiple caches
 - A write by one processor may not become visible
 - Accessing stale value in their caches
 - *Cache coherence problem*
- What do we do about it?
 - Organize the mem hierarchy to make it go away
 - Detect and take actions to eliminate the problem

Architectural Building Blocks

- Bus Transactions
 - fundamental system design abstraction
 - single set of wires connect several devices
 - bus protocol: arbitration, command/addr, data
 - Every device observes every transaction
- Cache block state transition diagram
 - FSM specifying how disposition of block changes
 - invalid, valid, dirty


Design Choices

- Cache Controller
 - Updates state of blocks in response to processor and snoop events
 - Generates bus transactions
- Action Choices
 - Write-through vs. Write-back
 - Invalidate vs. Update



Multi-core CPU chip

- All multiprocessor concepts apply
- The cores fit on a single processor socket
- Also called CMP (Chip Multi-Processor)

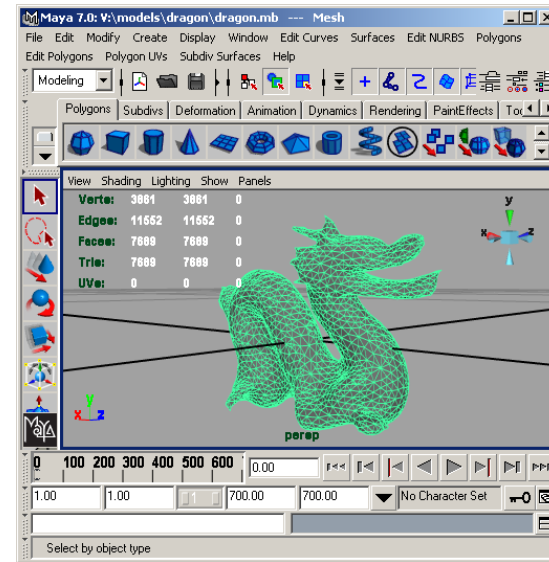


Why Multiple Processors

- Moore's Law in CMOS
- Deeply Pipelined circuits
 - Large size processors
 - Difficult hardware design and verification
 - Large design teams necessary
- Many Dumb But Useful Applications
 - Many new applications are multithreaded
 - General trend in computer architecture
- Marketing Strategy
 - Increased number of processors
 - Potential for instantaneous high performance

Today's Cloud Application

- Database servers
- Web servers (Web commerce)
- Compilers
- Multimedia applications
- Scientific applications, CAD/CAM
- In general, applications with *Thread+ level parallelism* (as opposed to instruction-level parallelism)



Each can run on its own core

