

# CSCI567 Machine Learning (Spring 2023)

## Week 1: Course overview, KNN, ML system

Prof. Yan Liu

University of Southern California

# Outline

- 1 Course Overview
- 2 Overview of machine learning
- 3 Nearest Neighbor Classifier (NNC)
- 4 Some theory on NNC

# Outline

- 1 Course Overview
- 2 Overview of machine learning
- 3 Nearest Neighbor Classifier (NNC)
- 4 Some theory on NNC

# Outline

- 1 Course Overview
- 2 Overview of machine learning
- 3 Nearest Neighbor Classifier (NNC)
- 4 Some theory on NNC


# What is machine learning?

## One possible definition<sup>1</sup>

a set of methods that can automatically *detect patterns* in data, and then use the uncovered patterns to *predict future data*, or to perform other kinds of decision making *under uncertainty*


# Example: detect patterns

How the temperature has been changing?



## Example: detect patterns

How the temperature has been changing?




### Patterns

- Seems going up
- Repeated periods of going up and down.

# How do we describe the pattern?

## Build a model: fit the data with a polynomial function




- The model is not accurate for individual years
- But collectively, the model captures the major trend
- Still, not able to model the pattern of the *repeated up and down*



# Predicting future

## What is temperature of 2010?



- Again, the model is not accurate for that specific year
- But then, it is close to the actual one

# What we have learned from this example?

## Key ingredients in machine learning

- Data  
collected from past observation (we often call them *training data*)

# What we have learned from this example?

## Key ingredients in machine learning

- Data  
collected from past observation (we often call them *training data*)
- Modeling  
devised to capture the patterns in the data
  - The model does not have to be true — “All models are wrong, but some are useful” by George Box.

# What we have learned from this example?

## Key ingredients in machine learning

- Data  
collected from past observation (we often call them *training data*)
- Modeling  
designed to capture the patterns in the data
  - The model does not have to be true — “All models are wrong, but some are useful” by George Box.
- Prediction  
apply the model to forecast what is going to happen in future

# A rich history of applying statistical learning methods


## Recognizing flowers (by R. Fisher, 1936)

Types of Iris: setosa, versicolor, and virginica




# Huge success 20 years ago

## Recognizing handwritten zipcodes (AT&T Labs, late 1990s)




# More modern ones, in your social life

## Recognizing your friends on Facebook



# It might be possible to know about you than yourself

## Recommending what you might like





# Why is machine learning so hot?

- **Tons of consumer applications:**
  - speech recognition, information retrieval and search, email and document classification, stock price prediction, object recognition, biometrics, etc
  - Highly desirable expertise from industry: Google, Facebook, Microsoft, Uber, Twitter, IBM, LinkedIn, Amazon, ...

# Why is machine learning so hot?

- **Tons of consumer applications:**

- speech recognition, information retrieval and search, email and document classification, stock price prediction, object recognition, biometrics, etc
- Highly desirable expertise from industry: Google, Facebook, Microsoft, Uber, Twitter, IBM, LinkedIn, Amazon, ...

- **Enable scientific breakthrough**

- Climate science: understand global warming cause and effect
- Biology and genetics: identify disease-causing genes and gene networks
- Social science: social network analysis; social media analysis
- Business and finance: marketing, operation research
- Emerging ones: healthcare, energy, ...

# What is in machine learning?

## Different flavors of learning problems

- Supervised learning  
Aim to predict (as in previous examples)

# What is in machine learning?

## Different flavors of learning problems

- Supervised learning  
Aim to predict (as in previous examples)
- Unsupervised learning  
Aim to discover hidden and latent patterns and explore data

# What is in machine learning?

## Different flavors of learning problems

- Supervised learning  
Aim to predict (as in previous examples)
- Unsupervised learning  
Aim to discover hidden and latent patterns and explore data
- Reinforcement learning  
Aim to act optimally under uncertainty

# What is in machine learning?

## Different flavors of learning problems

- Supervised learning  
Aim to predict (as in previous examples)
- Unsupervised learning  
Aim to discover hidden and latent patterns and explore data
- Reinforcement learning  
Aim to act optimally under uncertainty
- Many other paradigms

# What is in machine learning?

## Different flavors of learning problems

- Supervised learning  
Aim to predict (as in previous examples)
- Unsupervised learning  
Aim to discover hidden and latent patterns and explore data
- Reinforcement learning  
Aim to act optimally under uncertainty
- Many other paradigms

## The focus and goal of this course

- Supervised learning (before midterm)
- Unsupervised learning (after midterm)

# Outline

- 1 Course Overview
- 2 Overview of machine learning
- 3 Nearest Neighbor Classifier (NNC)
  - Intuitive example
  - General setup for classification
  - Algorithm
  - How to measure performance
  - Variants, Parameters, and Tuning
  - Summary
- 4 Some theory on NNC




# Recognizing flowers

Types of Iris: *setosa*, *versicolor*, and *virginica*



# Measuring the properties of the flowers

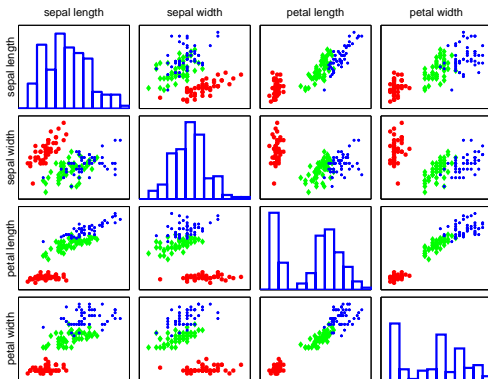
**Features and attributes: the widths and lengths of sepal and petal**



# Pairwise scatter plots of 131 flower specimens


Visualization of data helps identify the right learning model to use

Each colored point is a flower specimen: **setosa**, **versicolor**, **virginica**



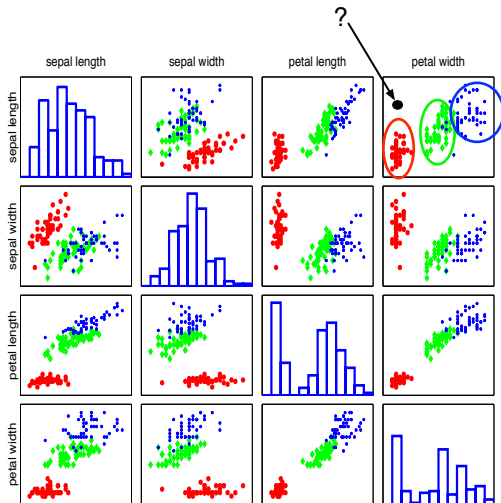
# Different types seem well-clustered and separable

Using two features: petal width and sepal length



# Labeling an unknown flower type

Closer to red cluster: so labeling it as **setosa**



# General setup for multi-class classification

## Training data (set)

- N samples/instances:  $\mathcal{D}^{\text{TRAIN}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)\}$

# General setup for multi-class classification

## Training data (set)

- N samples/instances:  $\mathcal{D}^{\text{TRAIN}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)\}$
- Each  $\mathbf{x}_n \in \mathbb{R}^D$  is called a feature vector.

# General setup for multi-class classification

## Training data (set)

- N samples/instances:  $\mathcal{D}^{\text{TRAIN}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)\}$
- Each  $\mathbf{x}_n \in \mathbb{R}^D$  is called a feature vector.
- Each  $y_n \in [C] = \{1, 2, \dots, C\}$  is called a label/class/category.



# General setup for multi-class classification

## Training data (set)

- N samples/instances:  $\mathcal{D}^{\text{TRAIN}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)\}$
- Each  $\mathbf{x}_n \in \mathbb{R}^D$  is called a feature vector.
- Each  $y_n \in [C] = \{1, 2, \dots, C\}$  is called a label/class/category.
- They are used for learning  $f : \mathbb{R}^D \rightarrow [C]$  for future prediction.

# General setup for multi-class classification

## Training data (set)

- N samples/instances:  $\mathcal{D}^{\text{TRAIN}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)\}$
- Each  $\mathbf{x}_n \in \mathbb{R}^D$  is called a feature vector.
- Each  $y_n \in [C] = \{1, 2, \dots, C\}$  is called a label/class/category.
- They are used for learning  $f : \mathbb{R}^D \rightarrow [C]$  for future prediction.

## Special case: binary classification

- Number of classes:  $C = 2$
- Conventional labels:  $\{0, 1\}$  or  $\{-1, +1\}$

# Often, data is conveniently organized as a table

## Ex: Iris data ([click here for all data](#))

- 4 features
- 3 classes

Fisher's *Iris* Data

Sepal length ↕	Sepal width ↕	Petal length ↕	Petal width ↕	Species ↕
5.1	3.5	1.4	0.2	<i>I. setosa</i>
4.9	3.0	1.4	0.2	<i>I. setosa</i>
4.7	3.2	1.3	0.2	<i>I. setosa</i>
4.6	3.1	1.5	0.2	<i>I. setosa</i>
5.0	3.6	1.4	0.2	<i>I. setosa</i>
5.4	3.9	1.7	0.4	<i>I. setosa</i>
4.6	3.4	1.4	0.3	<i>I. setosa</i>
5.0	3.4	1.5	0.2	<i>I. setosa</i>
4.4	2.9	1.4	0.2	<i>I. setosa</i>
4.9	3.1	1.5	0.1	<i>I. setosa</i>

# Nearest neighbor classification (NNC)

## Nearest neighbor

$$\mathbf{x}(1) = \mathbf{x}_{\text{nn}(\mathbf{x})}$$

where  $\text{nn}(\mathbf{x}) \in [\mathbf{N}] = \{1, 2, \dots, \mathbf{N}\}$ , i.e., the index to one of the training instances,

$$\text{nn}(\mathbf{x}) = \underset{n \in [\mathbf{N}]}{\operatorname{argmin}} \|\mathbf{x} - \mathbf{x}_n\|_2 = \underset{n \in [\mathbf{N}]}{\operatorname{argmin}} \sqrt{\sum_{d=1}^D (x_d - x_{nd})^2}$$

where  $\|\cdot\|_2$  is the  $L_2$ /Euclidean distance.

# Nearest neighbor classification (NNC)

## Nearest neighbor

$$\mathbf{x}(1) = \mathbf{x}_{\text{nn}(\mathbf{x})}$$

where  $\text{nn}(\mathbf{x}) \in [\mathbf{N}] = \{1, 2, \dots, N\}$ , i.e., the index to one of the training instances,

$$\text{nn}(\mathbf{x}) = \underset{n \in [\mathbf{N}]}{\operatorname{argmin}} \|\mathbf{x} - \mathbf{x}_n\|_2 = \underset{n \in [\mathbf{N}]}{\operatorname{argmin}} \sqrt{\sum_{d=1}^D (x_d - x_{nd})^2}$$


where  $\|\cdot\|_2$  is the  $L_2$ /Euclidean distance.

## Classification rule

$$y = f(\mathbf{x}) = y_{\text{nn}(\mathbf{x})}$$

## Visual example

In this 2-dimensional example, the nearest point to  $x$  is a **red training instance**, thus,  $x$  will be labeled as **red**.



## Example: classify Iris with two features

### Training data

ID (n)	petal width ( $x_1$ )	sepal length ( $x_2$ )	category ( $y$ )
1	0.2	5.1	setoas
2	1.4	7.0	versicolor
3	2.5	6.7	virginica
$\vdots$	$\vdots$	$\vdots$	

## Example: classify Iris with two features

### Training data

ID (n)	petal width ( $x_1$ )	sepal length ( $x_2$ )	category ( $y$ )
1	0.2	5.1	setoas
2	1.4	7.0	versicolor
3	2.5	6.7	virginica
$\vdots$	$\vdots$	$\vdots$	

### Flower with unknown category

petal width = 1.8 and sepal width = 6.4 (i.e.  $\mathbf{x} = (1.8, 6.4)$ )

Calculating distance  $\|\mathbf{x} - \mathbf{x}_n\|_2 = \sqrt{(x_1 - x_{n1})^2 + (x_2 - x_{n2})^2}$


ID	distance
1	1.75
2	0.72
3	0.76

Thus, the category is *versicolor*.



# Decision boundary

For every point in the space, we can determine its label using the NNC rule. This gives rise to a *decision boundary* that partitions the space into different regions.



# Is NNC doing the right thing for us?

## Intuition

We should compute **accuracy** — the percentage of data points being correctly classified, or the **error rate** — the percentage of data points being incorrectly classified. (accuracy + error rate = 1)

# Is NNC doing the right thing for us?

## Intuition

We should compute **accuracy** — the percentage of data points being correctly classified, or the **error rate** — the percentage of data points being incorrectly classified. (accuracy + error rate = 1)

## Defined on the training data set

$$A^{\text{TRAIN}} = \frac{1}{N} \sum_n \mathbb{I}[f(\mathbf{x}_n) == y_n], \quad \varepsilon^{\text{TRAIN}} = \frac{1}{N} \sum_n \mathbb{I}[f(\mathbf{x}_n) \neq y_n]$$

where  $\mathbb{I}[\cdot]$  is the indicator function.

# Is NNC doing the right thing for us?

## Intuition

We should compute **accuracy** — the percentage of data points being correctly classified, or the **error rate** — the percentage of data points being incorrectly classified. (accuracy + error rate = 1)

## Defined on the training data set


$$A^{\text{TRAIN}} = \frac{1}{N} \sum_n \mathbb{I}[f(\mathbf{x}_n) == y_n], \quad \varepsilon^{\text{TRAIN}} = \frac{1}{N} \sum_n \mathbb{I}[f(\mathbf{x}_n) \neq y_n]$$

where  $\mathbb{I}[\cdot]$  is the indicator function.

*Is this the right measure?*

# Example


Training data



What are  $A^{\text{TRAIN}}$  and  $\epsilon^{\text{TRAIN}}$ ?

# Example

Training data



What are  $A^{\text{TRAIN}}$  and  $\varepsilon^{\text{TRAIN}}$ ?

$$A^{\text{TRAIN}} = 100\%, \quad \varepsilon^{\text{TRAIN}} = 0\%$$

*For every training data point, its nearest neighbor is itself.*

# Test Error

Does it mean nearest neighbor is a very good algorithm?

# Test Error

Does it mean nearest neighbor is a very good algorithm?

*Not really, having zero training error is simple!*



# Test Error

Does it mean nearest neighbor is a very good algorithm?

*Not really, having zero training error is simple!*

We should care about accuracy when predicting unseen data

# Test Error

Does it mean nearest neighbor is a very good algorithm?

*Not really, having zero training error is simple!*

We should care about accuracy when predicting unseen data

## Test/Evaluation data

- $\mathcal{D}^{\text{TEST}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_M, y_M)\}$
- A fresh dataset, *not* overlap with training set.

# Test Error

Does it mean nearest neighbor is a very good algorithm?

*Not really, having zero training error is simple!*

We should care about accuracy when predicting unseen data

## Test/Evaluation data

- $\mathcal{D}^{\text{TEST}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_M, y_M)\}$
- A fresh dataset, *not* overlap with training set.
- Test accuracy and test error

$$A^{\text{TEST}} = \frac{1}{M} \sum_m \mathbb{I}[f(\mathbf{x}_m) == y_m], \quad \varepsilon^{\text{TEST}} = \frac{1}{M} \sum_M \mathbb{I}[f(\mathbf{x}_m) \neq y_m]$$

# Test Error

Does it mean nearest neighbor is a very good algorithm?

*Not really, having zero training error is simple!*

We should care about accuracy when predicting unseen data

## Test/Evaluation data

- $\mathcal{D}^{\text{TEST}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_M, y_M)\}$
- A fresh dataset, *not* overlap with training set.
- Test accuracy and test error

$$A^{\text{TEST}} = \frac{1}{M} \sum_m \mathbb{I}[f(\mathbf{x}_m) == y_m], \quad \varepsilon^{\text{TEST}} = \frac{1}{M} \sum_m \mathbb{I}[f(\mathbf{x}_m) \neq y_m]$$

- Good measurement of a classifier's performance

# Variant 1: measure nearness with other distances

Previously, we use the Euclidean distance

$$\text{nn}(\mathbf{x}) = \underset{n \in [\mathbf{N}]}{\operatorname{argmin}} \|\mathbf{x} - \mathbf{x}_n\|_2$$

## Variant 1: measure nearness with other distances


Previously, we use the Euclidean distance

$$\text{nn}(\mathbf{x}) = \underset{n \in [\mathbf{N}]}{\operatorname{argmin}} \|\mathbf{x} - \mathbf{x}_n\|_2$$

**Many other alternative distances**

E.g., the following  $L_1$  distance (i.e., city block distance, or Manhattan distance)

$$\|\mathbf{x} - \mathbf{x}_n\|_1 = \sum_{d=1}^D |x_d - x_{nd}|$$



Green line is Euclidean distance.  
Red, Blue, and Yellow lines are  $L_1$  distance

# Variant 1: measure nearness with other distances

Previously, we use the Euclidean distance

$$\text{nn}(\mathbf{x}) = \underset{n \in [\mathbf{N}]}{\operatorname{argmin}} \|\mathbf{x} - \mathbf{x}_n\|_2$$


**Many other alternative distances**

E.g., the following  $L_1$  distance (i.e., city block distance, or Manhattan distance)

$$\|\mathbf{x} - \mathbf{x}_n\|_1 = \sum_{d=1}^D |x_d - x_{nd}|$$

More generally,  $L_p$  distance (for  $p \geq 1$ ):

$$\|\mathbf{x} - \mathbf{x}_n\|_p = \left( \sum_d |x_d - x_{nd}|^p \right)^{1/p}$$



Green line is Euclidean distance.  
Red, Blue, and Yellow lines are  $L_1$  distance

## Variant 2: K-nearest neighbor (KNN)

### Increase the number of nearest neighbors to use?

- 1-nearest neighbor:  $\text{nn}_1(\mathbf{x}) = \operatorname{argmin}_{n \in [\mathbf{N}]} \|\mathbf{x} - \mathbf{x}_n\|_2$
- 2-nearest neighbor:  $\text{nn}_2(\mathbf{x}) = \operatorname{argmin}_{n \in [\mathbf{N}] - \text{nn}_1(\mathbf{x})} \|\mathbf{x} - \mathbf{x}_n\|_2$
- 3-nearest neighbor:  $\text{nn}_3(\mathbf{x}) = \operatorname{argmin}_{n \in [\mathbf{N}] - \text{nn}_1(\mathbf{x}) - \text{nn}_2(\mathbf{x})} \|\mathbf{x} - \mathbf{x}_n\|_2$



## Variant 2: K-nearest neighbor (KNN)

### Increase the number of nearest neighbors to use?

- 1-nearest neighbor:  $\text{nn}_1(\mathbf{x}) = \operatorname{argmin}_{n \in [\mathbf{N}]} \|\mathbf{x} - \mathbf{x}_n\|_2$
- 2-nearest neighbor:  $\text{nn}_2(\mathbf{x}) = \operatorname{argmin}_{n \in [\mathbf{N}] - \text{nn}_1(\mathbf{x})} \|\mathbf{x} - \mathbf{x}_n\|_2$
- 3-nearest neighbor:  $\text{nn}_3(\mathbf{x}) = \operatorname{argmin}_{n \in [\mathbf{N}] - \text{nn}_1(\mathbf{x}) - \text{nn}_2(\mathbf{x})} \|\mathbf{x} - \mathbf{x}_n\|_2$

### The set of K-nearest neighbor

$$\text{knn}(\mathbf{x}) = \{\text{nn}_1(\mathbf{x}), \text{nn}_2(\mathbf{x}), \dots, \text{nn}_K(\mathbf{x})\}$$

## Variant 2: K-nearest neighbor (KNN)

### Increase the number of nearest neighbors to use?

- 1-nearest neighbor:  $\text{nn}_1(\mathbf{x}) = \operatorname{argmin}_{n \in [\mathbf{N}]} \|\mathbf{x} - \mathbf{x}_n\|_2$
- 2-nearest neighbor:  $\text{nn}_2(\mathbf{x}) = \operatorname{argmin}_{n \in [\mathbf{N}] - \text{nn}_1(\mathbf{x})} \|\mathbf{x} - \mathbf{x}_n\|_2$
- 3-nearest neighbor:  $\text{nn}_3(\mathbf{x}) = \operatorname{argmin}_{n \in [\mathbf{N}] - \text{nn}_1(\mathbf{x}) - \text{nn}_2(\mathbf{x})} \|\mathbf{x} - \mathbf{x}_n\|_2$

### The set of K-nearest neighbor

$$\text{knn}(\mathbf{x}) = \{\text{nn}_1(\mathbf{x}), \text{nn}_2(\mathbf{x}), \dots, \text{nn}_K(\mathbf{x})\}$$

Note: with  $\mathbf{x}(k) = \mathbf{x}_{\text{nn}_k(\mathbf{x})}$ , we have

$$\|\mathbf{x} - \mathbf{x}(1)\|_2^2 \leq \|\mathbf{x} - \mathbf{x}(2)\|_2^2 \leq \dots \leq \|\mathbf{x} - \mathbf{x}(K)\|_2^2$$

# How to classify with $K$ neighbors?

## Classification rule

- Every neighbor votes: naturally  $x_n$  votes for its label  $y_n$ .

# How to classify with $K$ neighbors?

## Classification rule

- Every neighbor votes: naturally  $\mathbf{x}_n$  votes for its label  $y_n$ .
- Aggregate everyone's vote on a class label  $c$

$$v_c = \sum_{n \in \text{knn}(\mathbf{x})} \mathbb{I}(y_n == c), \quad \forall \quad c \in [\mathbf{C}]$$

# How to classify with $K$ neighbors?

## Classification rule




- Every neighbor votes: naturally  $\mathbf{x}_n$  votes for its label  $y_n$ .
- Aggregate everyone's vote on a class label  $c$

$$v_c = \sum_{n \in \text{knn}(\mathbf{x})} \mathbb{I}(y_n == c), \quad \forall \quad c \in [\mathbf{C}]$$


- Predict with the majority

$$y = f(\mathbf{x}) = \operatorname{argmax}_{c \in [\mathbf{C}]} v_c$$

# Example




# Decision boundary



When  $K$  increases, the decision boundary becomes smoother.

# Decision boundary



When  $K$  increases, the decision boundary becomes smoother.

*What happens when  $K = N$ ?*



## Variant 3: Preprocessing data

One issue of NNC: *distances depend on units of the features!*

## Variant 3: Preprocessing data

One issue of NNC: *distances depend on units of the features!*

**One solution: preprocess data so it looks more “normalized”.**

## Variant 3: Preprocessing data

One issue of NNC: *distances depend on units of the features!*

**One solution: preprocess data so it looks more “normalized”.**

Example:

- compute the means and standard deviations in each feature

$$\bar{x}_d = \frac{1}{N} \sum_n x_{nd}, \quad s_d^2 = \frac{1}{N-1} \sum_n (x_{nd} - \bar{x}_d)^2$$

- Scale the feature accordingly

$$x_{nd} \leftarrow \frac{x_{nd} - \bar{x}_d}{s_d}$$

## Variant 3: Preprocessing data

One issue of NNC: *distances depend on units of the features!*

**One solution: preprocess data so it looks more “normalized”.**

Example:

- compute the means and standard deviations in each feature

$$\bar{x}_d = \frac{1}{N} \sum_n x_{nd}, \quad s_d^2 = \frac{1}{N-1} \sum_n (x_{nd} - \bar{x}_d)^2$$

- Scale the feature accordingly

$$x_{nd} \leftarrow \frac{x_{nd} - \bar{x}_d}{s_d}$$

Many other ways of normalizing data.

# Which variants should we use?

## Hyper-parameters in NNC

- The distance measure (e.g. the parameter  $p$  for  $L_p$  norm)
- $K$  (i.e. how many nearest neighbor?)
- Different ways of preprocessing

# Which variants should we use?

## Hyper-parameters in NNC

- The distance measure (e.g. the parameter  $p$  for  $L_p$  norm)
- $K$  (i.e. how many nearest neighbor?)
- Different ways of preprocessing

*Most algorithms have hyper-parameters. Tuning them is a significant part of applying an algorithm.*

# Tuning via a development dataset

## Training data

- N samples/instances:  $\mathcal{D}^{\text{TRAIN}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)\}$
- They are used for learning  $f(\cdot)$

## Test data

- M samples/instances:  $\mathcal{D}^{\text{TEST}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_M, y_M)\}$
- They are used for assessing how well  $f(\cdot)$  will do.

# Tuning via a development dataset

## Training data

- N samples/instances:  $\mathcal{D}^{\text{TRAIN}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)\}$
- They are used for learning  $f(\cdot)$

## Test data

- M samples/instances:  $\mathcal{D}^{\text{TEST}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_M, y_M)\}$
- They are used for assessing how well  $f(\cdot)$  will do.

## Development/Validation data

- L samples/instances:  $\mathcal{D}^{\text{DEV}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_L, y_L)\}$
- They are used to optimize hyper-parameter(s).



# Tuning via a development dataset

## Training data

- N samples/instances:  $\mathcal{D}^{\text{TRAIN}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)\}$
- They are used for learning  $f(\cdot)$

## Test data

- M samples/instances:  $\mathcal{D}^{\text{TEST}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_M, y_M)\}$
- They are used for assessing how well  $f(\cdot)$  will do.

## Development/Validation data

- L samples/instances:  $\mathcal{D}^{\text{DEV}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_L, y_L)\}$
- They are used to optimize hyper-parameter(s).

These three sets should *not* overlap!

# Recipe

- For each possible value of the hyperparameter (e.g.  $K = 1, 3, \dots$ )
  - Train a model using  $\mathcal{D}^{\text{TRAIN}}$
  - Evaluate the performance of the model on  $\mathcal{D}^{\text{DEV}}$

# Recipe

- For each possible value of the hyperparameter (e.g.  $K = 1, 3, \dots$ )
  - Train a model using  $\mathcal{D}^{\text{TRAIN}}$
  - Evaluate the performance of the model on  $\mathcal{D}^{\text{DEV}}$
- Choose the model with the best performance on  $\mathcal{D}^{\text{DEV}}$

# Recipe


- For each possible value of the hyperparameter (e.g.  $K = 1, 3, \dots$ )
  - Train a model using  $\mathcal{D}^{\text{TRAIN}}$
  - Evaluate the performance of the model on  $\mathcal{D}^{\text{DEV}}$
- Choose the model with the best performance on  $\mathcal{D}^{\text{DEV}}$
- Evaluate the model on  $\mathcal{D}^{\text{TEST}}$

# S-fold Cross-validation

## What if we do not have a development set?

- Split the training data into  $S$  equal parts.

$S = 5$ : 5-fold cross validation




# S-fold Cross-validation

## What if we do not have a development set?

- Split the training data into  $S$  equal parts.
- Use each part *in turn* as a development dataset and use the others as a training dataset.

$S = 5$ : 5-fold cross validation




# S-fold Cross-validation

## What if we do not have a development set?

- Split the training data into  $S$  equal parts.
- Use each part *in turn* as a development dataset and use the others as a training dataset.
- Choose the hyper-parameter leading to best *average* performance.

$S = 5$ : 5-fold cross validation




# S-fold Cross-validation

## What if we do not have a development set?

- Split the training data into  $S$  equal parts.
- Use each part *in turn* as a development dataset and use the others as a training dataset.
- Choose the hyper-parameter leading to best *average* performance.

$S = 5$ : 5-fold cross validation



*Special case:*  $S = N$ , called leave-one-out.



# Cross-validation recipe

- Split the training data into  $S$  equal parts. Denote each part as  $\mathcal{D}_s^{\text{TRAIN}}$ .

# Cross-validation recipe

- Split the training data into  $S$  equal parts. Denote each part as  $\mathcal{D}_s^{\text{TRAIN}}$ .
- For each possible value of the hyper-parameter (e.g.  $K = 1, 3, \dots$ )
  - For every  $s \in [S]$ 
    - Train a model using  $\mathcal{D}_{\setminus s}^{\text{TRAIN}} = \mathcal{D}^{\text{TRAIN}} - \mathcal{D}_s^{\text{TRAIN}}$
    - Evaluate the performance of the model on  $\mathcal{D}_s^{\text{TRAIN}}$
  - Average the  $S$  performance metrics

# Cross-validation recipe

- Split the training data into  $S$  equal parts. Denote each part as  $\mathcal{D}_s^{\text{TRAIN}}$ .
- For each possible value of the hyper-parameter (e.g.  $K = 1, 3, \dots$ )
  - For every  $s \in [S]$ 
    - Train a model using  $\mathcal{D}_{\setminus s}^{\text{TRAIN}} = \mathcal{D}^{\text{TRAIN}} - \mathcal{D}_s^{\text{TRAIN}}$
    - Evaluate the performance of the model on  $\mathcal{D}_s^{\text{TRAIN}}$
  - Average the  $S$  performance metrics
- Choose the hyper-parameter with the best averaged performance

# Cross-validation recipe

- Split the training data into  $S$  equal parts. Denote each part as  $\mathcal{D}_s^{\text{TRAIN}}$ .
- For each possible value of the hyper-parameter (e.g.  $K = 1, 3, \dots$ )
  - For every  $s \in [S]$ 
    - Train a model using  $\mathcal{D}_{\setminus s}^{\text{TRAIN}} = \mathcal{D}^{\text{TRAIN}} - \mathcal{D}_s^{\text{TRAIN}}$
    - Evaluate the performance of the model on  $\mathcal{D}_s^{\text{TRAIN}}$
  - Average the  $S$  performance metrics
- Choose the hyper-parameter with the best averaged performance
- **Use the best hyper-parameter to train a model using all  $\mathcal{D}^{\text{train}}$**

# Cross-validation recipe

- Split the training data into  $S$  equal parts. Denote each part as  $\mathcal{D}_s^{\text{TRAIN}}$ .
- For each possible value of the hyper-parameter (e.g.  $K = 1, 3, \dots$ )
  - For every  $s \in [S]$ 
    - Train a model using  $\mathcal{D}_{\setminus s}^{\text{TRAIN}} = \mathcal{D}^{\text{TRAIN}} - \mathcal{D}_s^{\text{TRAIN}}$
    - Evaluate the performance of the model on  $\mathcal{D}_s^{\text{TRAIN}}$
  - Average the  $S$  performance metrics
- Choose the hyper-parameter with the best averaged performance
- **Use the best hyper-parameter to train a model using all  $\mathcal{D}^{\text{train}}$**
- Evaluate the model on  $\mathcal{D}^{\text{TEST}}$

# Mini-summary

## Advantages of NNC

- Simple, easy to implement (wildly used in practice)

# Mini-summary

## Advantages of NNC

- Simple, easy to implement (wildly used in practice)

## Disadvantages of NNC

- Computationally intensive for large-scale problems:  $O(ND)$  for each prediction *naively*.

# Mini-summary

## Advantages of NNC

- Simple, easy to implement (wildly used in practice)

## Disadvantages of NNC

- Computationally intensive for large-scale problems:  $O(ND)$  for each prediction *naively*.
- Need to “*carry*” the training data around. This type of method is called *nonparametric*.



# Mini-summary

## Advantages of NNC

- Simple, easy to implement (wildly used in practice)

## Disadvantages of NNC

- Computationally intensive for large-scale problems:  $O(ND)$  for each prediction *naively*.
- Need to “*carry*” the training data around. This type of method is called *nonparametric*.
- Choosing the right hyper-parameters can be involved.

# Mini-summary

**Typical steps** of developing a machine learning system:

- Collect data, split into training, development, and test sets.
- Train a model with a machine learning algorithm. Most often we apply cross-validation to tune hyper-parameters.
- Evaluate using the test data and report performance.
- Use the model to predict future/make decisions.

# Outline

- 1 Course Overview
- 2 Overview of machine learning
- 3 Nearest Neighbor Classifier (NNC)
- 4 Some theory on NNC
  - Step 1: Expected risk
  - Step 2: The ideal classifier
  - Step 3: Comparing NNC to the ideal classifier

# How good is NNC really?

To answer this question, we proceed in 3 steps

# How good is NNC really?

To answer this question, we proceed in 3 steps

- 1 Define *more carefully* a performance metric for a classifier.

# How good is NNC really?

To answer this question, we proceed in 3 steps

- 1 Define *more carefully* a performance metric for a classifier.
- 2 Hypothesize an ideal classifier - *the best possible one*.

# How good is NNC really?

To answer this question, we proceed in 3 steps

- 1 Define *more carefully* a performance metric for a classifier.
- 2 Hypothesize an ideal classifier - *the best possible one*.
- 3 Compare NNC to the ideal one.

# Why does test error make sense?

Test error makes sense only when training set and test set are correlated.



# Why does test error make sense?

Test error makes sense only when training set and test set are correlated.

**Most standard assumption:** every data point  $(x, y)$  (from  $\mathcal{D}^{\text{TRAIN}}$ ,  $\mathcal{D}^{\text{DEV}}$ , or  $\mathcal{D}^{\text{TEST}}$ ) is an *independent and identically distributed (i.i.d.)* sample of an unknown joint distribution  $\mathcal{P}$ .

- often written as  $(x, y) \stackrel{i.i.d.}{\sim} \mathcal{P}$

# Why does test error make sense?

Test error makes sense only when training set and test set are correlated.

**Most standard assumption:** every data point  $(x, y)$  (from  $\mathcal{D}^{\text{TRAIN}}$ ,  $\mathcal{D}^{\text{DEV}}$ , or  $\mathcal{D}^{\text{TEST}}$ ) is an *independent and identically distributed (i.i.d.)* sample of an unknown joint distribution  $\mathcal{P}$ .

- often written as  $(x, y) \stackrel{i.i.d.}{\sim} \mathcal{P}$

Test error of a fixed classifier is therefore a *random variable*.

# Why does test error make sense?

Test error makes sense only when training set and test set are correlated.

**Most standard assumption:** every data point  $(x, y)$  (from  $\mathcal{D}^{\text{TRAIN}}$ ,  $\mathcal{D}^{\text{DEV}}$ , or  $\mathcal{D}^{\text{TEST}}$ ) is an *independent and identically distributed (i.i.d.)* sample of an unknown joint distribution  $\mathcal{P}$ .

- often written as  $(x, y) \stackrel{i.i.d.}{\sim} \mathcal{P}$

Test error of a fixed classifier is therefore a *random variable*.

Need a more “certain” measure of performance (so it’s easy to compare different classifiers for example).

# Expected error

What about the **expectation** of this random variable?

$$\mathbb{E}[\epsilon^{\text{TEST}}]$$

# Expected error

What about the **expectation** of this random variable?

$$\mathbb{E}[\epsilon^{\text{TEST}}] = \frac{1}{M} \sum_{m=1}^M \mathbb{E}_{(\mathbf{x}_m, y_m) \sim \mathcal{P}} \mathbb{I}[f(\mathbf{x}_m) \neq y_m]$$

# Expected error

What about the **expectation** of this random variable?

$$\mathbb{E}[\epsilon^{\text{TEST}}] = \frac{1}{M} \sum_{m=1}^M \mathbb{E}_{(\mathbf{x}_m, y_m) \sim \mathcal{P}} \mathbb{I}[f(\mathbf{x}_m) \neq y_m] = \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{P}} \mathbb{I}[f(\mathbf{x}) \neq y]$$

# Expected error

What about the **expectation** of this random variable?

$$\mathbb{E}[\epsilon^{\text{TEST}}] = \frac{1}{M} \sum_{m=1}^M \mathbb{E}_{(\mathbf{x}_m, y_m) \sim \mathcal{P}} \mathbb{I}[f(\mathbf{x}_m) \neq y_m] = \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{P}} \mathbb{I}[f(\mathbf{x}) \neq y]$$

- i.e. the expected error/mistake of  $f$

# Expected error

What about the **expectation** of this random variable?

$$\mathbb{E}[\epsilon^{\text{TEST}}] = \frac{1}{M} \sum_{m=1}^M \mathbb{E}_{(\mathbf{x}_m, y_m) \sim \mathcal{P}} \mathbb{I}[f(\mathbf{x}_m) \neq y_m] = \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{P}} \mathbb{I}[f(\mathbf{x}) \neq y]$$

- i.e. the expected error/mistake of  $f$

Test error is a proxy of expected error. *The larger the test set, the better the approximation.*



# Expected error

What about the **expectation** of this random variable?

$$\mathbb{E}[\epsilon^{\text{TEST}}] = \frac{1}{M} \sum_{m=1}^M \mathbb{E}_{(\mathbf{x}_m, y_m) \sim \mathcal{P}} \mathbb{I}[f(\mathbf{x}_m) \neq y_m] = \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{P}} \mathbb{I}[f(\mathbf{x}) \neq y]$$

- i.e. the expected error/mistake of  $f$

Test error is a proxy of expected error. *The larger the test set, the better the approximation.*

What about the expectation of training error? Is training error a good proxy of expected error?

# Expected risk

More generally, for a loss function  $L(y', y)$ ,

- e.g.  $L(y', y) = \mathbb{I}[y' \neq y]$ , called *0-1 loss*.
- many more other losses as we will see.

# Expected risk

More generally, for a loss function  $L(y', y)$ ,

- e.g.  $L(y', y) = \mathbb{I}[y' \neq y]$ , called *0-1 loss*.
- many more other losses as we will see.

the *expected risk* of  $f$  is defined as

$$R(f) = \mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{P}} L(f(\mathbf{x}), y)$$

# Expected risk

More generally, for a loss function  $L(y', y)$ ,

- e.g.  $L(y', y) = \mathbb{I}[y' \neq y]$ , called *0-1 loss*. **Default**
- many more other losses as we will see.

the *expected risk* of  $f$  is defined as

$$R(f) = \mathbb{E}_{(x,y) \sim \mathcal{P}} L(f(x), y)$$

# Bayes optimal classifier

What should we predict for  $x$ , *knowing*  $\mathcal{P}(y|x)$ ?

# Bayes optimal classifier

What should we predict for  $x$ , *knowing*  $\mathcal{P}(y|x)$ ?

**Bayes optimal classifier:**  $f^*(x) = \operatorname{argmax}_{c \in [C]} \mathcal{P}(c|x)$ .

# Bayes optimal classifier

What should we predict for  $x$ , *knowing*  $\mathcal{P}(y|x)$ ?

**Bayes optimal classifier:**  $f^*(x) = \operatorname{argmax}_{c \in [C]} \mathcal{P}(c|x)$ .

**The optimal risk:**  $R(f^*) = \mathbb{E}_{x \sim \mathcal{P}_x} [1 - \max_{c \in [C]} \mathcal{P}(c|x)]$  where  $\mathcal{P}_x$  is the marginal distribution of  $x$ .

# Bayes optimal classifier

What should we predict for  $x$ , *knowing*  $\mathcal{P}(y|x)$ ?

**Bayes optimal classifier:**  $f^*(x) = \operatorname{argmax}_{c \in [C]} \mathcal{P}(c|x)$ .

**The optimal risk:**  $R(f^*) = \mathbb{E}_{x \sim \mathcal{P}_x} [1 - \max_{c \in [C]} \mathcal{P}(c|x)]$  where  $\mathcal{P}_x$  is the marginal distribution of  $x$ .

It is easy to show  $R(f^*) \leq R(f)$  for any  $f$ .



# Bayes optimal classifier

What should we predict for  $x$ , *knowing*  $\mathcal{P}(y|x)$ ?

**Bayes optimal classifier:**  $f^*(x) = \operatorname{argmax}_{c \in [C]} \mathcal{P}(c|x)$ .

**The optimal risk:**  $R(f^*) = \mathbb{E}_{x \sim \mathcal{P}_x} [1 - \max_{c \in [C]} \mathcal{P}(c|x)]$  where  $\mathcal{P}_x$  is the marginal distribution of  $x$ .

It is easy to show  $R(f^*) \leq R(f)$  for any  $f$ .

For special case  $C = 2$ , let  $\eta(x) = \mathcal{P}(0|x)$ , then

$$R(f^*) = \mathbb{E}_{x \sim \mathcal{P}_x} [\min\{\eta(x), 1 - \eta(x)\}].$$

# Comparing NNC to Bayes optimal classifier

**Come back to the question: how good is NNC?**

# Comparing NNC to Bayes optimal classifier

## Come back to the question: how good is NNC?

Theorem (Cover and Hart, 1967)

*Let  $f_N$  be the 1-nearest neighbor binary classifier using  $N$  training data points, we have (under mild conditions)*

$$R(f^*) \leq \lim_{N \rightarrow \infty} \mathbb{E}[R(f_N)] \leq 2R(f^*)$$

*i.e., expected risk of NNC in the limit is at most twice of the best possible.*

# Comparing NNC to Bayes optimal classifier

## Come back to the question: how good is NNC?

Theorem (Cover and Hart, 1967)

*Let  $f_N$  be the 1-nearest neighbor binary classifier using  $N$  training data points, we have (under mild conditions)*

$$R(f^*) \leq \lim_{N \rightarrow \infty} \mathbb{E}[R(f_N)] \leq 2R(f^*)$$

*i.e., expected risk of NNC in the limit is at most twice of the best possible.*

A pretty strong guarantee.

In particular,  $R(f^*) = 0$  implies  $\mathbb{E}[R(f_N)] \rightarrow 0$ .

# Proof sketch

**Fact:**  $x(1) \rightarrow x$  with probability 1

# Proof sketch

**Fact:**  $x(1) \rightarrow x$  with probability 1

$$\mathbb{E}[R(f_N)] = \mathbb{E}[\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{P}} \mathbb{I}[f_N(\mathbf{x}) \neq y]]$$

# Proof sketch

**Fact:**  $x(1) \rightarrow x$  with probability 1

$$\begin{aligned}\mathbb{E}[R(f_N)] &= \mathbb{E}[\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{P}} \mathbb{I}[f_N(\mathbf{x}) \neq y]] \\ &\rightarrow \mathbb{E}_{\mathbf{x} \sim \mathcal{P}_x} \mathbb{E}_{y, y' \stackrel{i.i.d.}{\sim} \mathcal{P}(\cdot | \mathbf{x})} [\mathbb{I}[y' \neq y]]\end{aligned}$$

# Proof sketch

**Fact:**  $x(1) \rightarrow x$  with probability 1

$$\begin{aligned}\mathbb{E}[R(f_N)] &= \mathbb{E}[\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{P}} \mathbb{I}[f_N(\mathbf{x}) \neq y]] \\ &\rightarrow \mathbb{E}_{\mathbf{x} \sim \mathcal{P}_x} \mathbb{E}_{y, y' \stackrel{i.i.d.}{\sim} \mathcal{P}(\cdot | \mathbf{x})} [\mathbb{I}[y' \neq y]] \\ &= \mathbb{E}_{\mathcal{P}_x} \mathbb{E}_{y, y' \stackrel{i.i.d.}{\sim} \mathcal{P}(\cdot | \mathbf{x})} [\mathbb{I}[y' = 0 \text{ and } y = 1] + \mathbb{I}[y' = 1 \text{ and } y = 0]]\end{aligned}$$



# Proof sketch

**Fact:**  $x(1) \rightarrow x$  with probability 1

$$\begin{aligned}\mathbb{E}[R(f_N)] &= \mathbb{E}[\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{P}} \mathbb{I}[f_N(\mathbf{x}) \neq y]] \\ &\rightarrow \mathbb{E}_{\mathbf{x} \sim \mathcal{P}_x} \mathbb{E}_{y, y' \stackrel{i.i.d.}{\sim} \mathcal{P}(\cdot | \mathbf{x})} [\mathbb{I}[y' \neq y]] \\ &= \mathbb{E}_{\mathcal{P}_x} \mathbb{E}_{y, y' \stackrel{i.i.d.}{\sim} \mathcal{P}(\cdot | \mathbf{x})} [\mathbb{I}[y' = 0 \text{ and } y = 1] + \mathbb{I}[y' = 1 \text{ and } y = 0]] \\ &= \mathbb{E}_{\mathcal{P}_x} [\eta(x)(1 - \eta(x)) + (1 - \eta(x))\eta(x)]\end{aligned}$$

# Proof sketch

**Fact:**  $x(1) \rightarrow x$  with probability 1

$$\begin{aligned}\mathbb{E}[R(f_N)] &= \mathbb{E}[\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{P}} \mathbb{I}[f_N(\mathbf{x}) \neq y]] \\ &\rightarrow \mathbb{E}_{\mathbf{x} \sim \mathcal{P}_x} \mathbb{E}_{y, y' \stackrel{i.i.d.}{\sim} \mathcal{P}(\cdot | \mathbf{x})} [\mathbb{I}[y' \neq y]] \\ &= \mathbb{E}_{\mathcal{P}_x} \mathbb{E}_{y, y' \stackrel{i.i.d.}{\sim} \mathcal{P}(\cdot | \mathbf{x})} [\mathbb{I}[y' = 0 \text{ and } y = 1] + \mathbb{I}[y' = 1 \text{ and } y = 0]] \\ &= \mathbb{E}_{\mathcal{P}_x} [\eta(x)(1 - \eta(x)) + (1 - \eta(x))\eta(x)] \\ &= 2\mathbb{E}_{\mathcal{P}_x} [\eta(x)(1 - \eta(x))]\end{aligned}$$

# Proof sketch

**Fact:**  $x(1) \rightarrow x$  with probability 1

$$\begin{aligned}\mathbb{E}[R(f_N)] &= \mathbb{E}[\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{P}} \mathbb{I}[f_N(\mathbf{x}) \neq y]] \\ &\rightarrow \mathbb{E}_{\mathbf{x} \sim \mathcal{P}_x} \mathbb{E}_{y, y' \stackrel{i.i.d.}{\sim} \mathcal{P}(\cdot | \mathbf{x})} [\mathbb{I}[y' \neq y]] \\ &= \mathbb{E}_{\mathcal{P}_x} \mathbb{E}_{y, y' \stackrel{i.i.d.}{\sim} \mathcal{P}(\cdot | \mathbf{x})} [\mathbb{I}[y' = 0 \text{ and } y = 1] + \mathbb{I}[y' = 1 \text{ and } y = 0]] \\ &= \mathbb{E}_{\mathcal{P}_x} [\eta(x)(1 - \eta(x)) + (1 - \eta(x))\eta(x)] \\ &= 2\mathbb{E}_{\mathcal{P}_x} [\eta(x)(1 - \eta(x))] \\ &\leq 2\mathbb{E}_{\mathcal{P}_x} [\min\{\eta(x), (1 - \eta(x))\}]\end{aligned}$$

# Proof sketch

**Fact:**  $x(1) \rightarrow x$  with probability 1

$$\begin{aligned}
 \mathbb{E}[R(f_N)] &= \mathbb{E}[\mathbb{E}_{(\mathbf{x}, y) \sim \mathcal{P}} \mathbb{I}[f_N(\mathbf{x}) \neq y]] \\
 &\rightarrow \mathbb{E}_{\mathbf{x} \sim \mathcal{P}_x} \mathbb{E}_{y, y' \stackrel{i.i.d.}{\sim} \mathcal{P}(\cdot | \mathbf{x})} [\mathbb{I}[y' \neq y]] \\
 &= \mathbb{E}_{\mathcal{P}_x} \mathbb{E}_{y, y' \stackrel{i.i.d.}{\sim} \mathcal{P}(\cdot | \mathbf{x})} [\mathbb{I}[y' = 0 \text{ and } y = 1] + \mathbb{I}[y' = 1 \text{ and } y = 0]] \\
 &= \mathbb{E}_{\mathcal{P}_x} [\eta(x)(1 - \eta(x)) + (1 - \eta(x))\eta(x)] \\
 &= 2\mathbb{E}_{\mathcal{P}_x} [\eta(x)(1 - \eta(x))] \\
 &\leq 2\mathbb{E}_{\mathcal{P}_x} [\min\{\eta(x), (1 - \eta(x))\}] \\
 &= 2R(f^*)
 \end{aligned}$$